

MITHUN KUMAR A

Business Analyst | Data & BI Analytics | BFSI Domain

Bangalore, India | +91 87783 14234 | mithunk98.0229@gmail.com

linkedin.com/in/mithunk98 | mithun-portfolio-flame.vercel.app | github.com/mithunk98

Professional Summary

Business Analyst with 2.5+ years across BFSI and retail, owning end-to-end delivery from requirements gathering through production automation for a lending book of **Rs. 253 Cr in cumulative disbursements**. Reports directly to the **Chief Risk Officer** and authors the monthly portfolio review presented to the **CEO, CRO and functional leadership** — translating C-suite questions into BRDs, data models, executive dashboards and decision-ready decks. Took recurring stakeholder reporting **from 0% to 100% automated**, shipping scheduled Python jobs into the company's production CI/CD platform, and currently leading development of **Lumen, an LLM-based credit underwriting agent**. Strong in SQL, Python, Tableau, Advanced Excel and executive PowerPoint storytelling, with experience designing controlled experiments to validate business interventions. Pursuing M.Sc. in Data Science.

Core Competencies

Business Analysis: Requirements Gathering, BRD/FRD, Functional Specification, Process Mapping, Gap Analysis, Stakeholder Management, JIRA

Analytics & BI: SQL (Advanced), Python, Tableau, Metabase, Advanced Excel, Executive PowerPoint Reporting, KPI Definition, Dashboard Design, Cohort Analysis

Experimentation & Statistics: A/B Test Design, Hypothesis Testing, Power Analysis & Sample Sizing, Stratified Randomisation, Confidence Intervals, Guardrail Metrics, Multiple-Comparison Correction, scipy / statsmodels

AI & LLM Systems: LLM-Based Decision Agents (Gemini), Model Evaluation & Calibration Frameworks, Blinded Benchmarking, Ground-Truth Leakage Prevention, Output Verification & Hallucination Detection

Data & Automation: ETL Pipeline Development, Reporting Automation, Automated Email Distribution, Scheduled Jobs (GitHub Actions CI/CD), Data Validation & Quality, PostgreSQL, MySQL, MariaDB, Pandas, NumPy, SQLAlchemy, openpyxl, AWS S3 & Athena, Metabase API, Git

Domain (BFSI): Credit & Collections Lifecycle, Portfolio Risk Reporting, Delinquency & Roll-Rate Forecasting, Vintage / Cohort Analysis, LTV, GNPA / NPA, Bucket Movement, Dealer & Agency Performance Management

Professional Experience

Vidyuttech Service Private Limited — Bangalore, India

EV lending & subscription financing; Rs. 253 Cr cumulative disbursements

Business Analyst

May 2025 – Present

Reporting directly to the Chief Risk Officer

- Own risk and portfolio analytics for a book of **Rs. 253 Cr in cumulative disbursements**, partnering with the CRO and business heads to define KPIs, gather requirements and deliver the reporting pack used in monthly leadership reviews.
- **Leading development of Lumen, an LLM-based credit underwriting agent** (Gemini) to automate approve / reject / refer decisions for EV financing — targeting **>90% accuracy on definitive decisions at 60%+ automation coverage** with a zero false-approval constraint; currently in calibration and benchmark evaluation.
- Built an **audit-grade evaluation framework** for the agent — blinded case-packet compilation, cryptographic hash-based decision freezing and air-gap verification to guarantee zero ground-truth leakage — and enforced an evidence bar requiring **multi-case corroboration** ($n \geq 3$) before accepting any calibration change to policy behaviour.
- Diagnosed **silent data-pipeline defects corrupting model inputs** through direct SQL investigation — join fan-out causing cross-record contamination, deduplication logic errors, missing offer-status coverage — and applied independent verification at every stage (PMT/FOIR recomputation, repeated-run stability testing, source-document cross-checks) to catch fabricated or duplicated data before it reached decisioning.
- **Designed and deployed the automated Risk Dashboard through the company's CI/CD pipeline** (Python, Metabase API, scheduled GitHub Actions workflow on the production jobs platform) — generating a **55-metric portfolio risk pack** across three product segments and two DPD definitions, auto-delivered to the CRO and risk leadership. Replaced a fully manual monthly build with a zero-touch, version-controlled pipeline.
- **Took partner risk reporting from 0% to 100% automated**, building a recipient-segmented distribution engine that emails each of **27+ dealership partners across 16 cities** only their own early-bounce and early-risk exposure, with area- and territory-manager CC lists resolved dynamically from source systems and rule-based suppression for excluded accounts.
- Author the **monthly Risk MFR**, the portfolio performance review presented to the **CEO, CRO and functional leadership** — covering GNPA cohort analysis, city-level flow-into-30+, manufacturer risk categorisation, NBFC-wise

bounce trends and subscription performance — and produce a tailored edition for the wider business audience.

- Automated **23+ daily and weekly operational reports** using Python and SQL, cutting manual reporting effort by **50%** (~20+ hours/week recovered) and eliminating recurring data-entry errors.
- Designed and delivered **15+ Metabase dashboards** covering portfolio health, collections performance, delinquency movement and agency KPIs across **4 partner agencies and 15+ cities**, reducing ad-hoc SQL requests from business teams by **40%**.
- Built **SQL-based ETL workflows across 10+ source tables** with automated validation checks, improving reporting accuracy by **30%** and enabling same-day stakeholder visibility.
- Developed roll-rate and delinquency forecasting models to project bucket movement and credit risk trends, improving repayment predictability by **25%** and informing collections resourcing decisions.
- Maintain a library of **75+ parameterised SQL and Python analytics scripts** spanning city, lender, manufacturer, portfolio and product-subscription views, enabling repeatable analysis without engineering support.
- Author requirement documents (BRD/FRD) and functional specifications for new reporting modules, tracking delivery through JIRA with business owners.

Business Analyst Trainee

Sep 2024 – May 2025

- Built the first standardised reporting layer for collections and portfolio operations, delivering **15+ dashboards** that replaced manually maintained spreadsheets.
- Created reusable SQL templates for LTV calculation, delinquency trending and agency-wise collection rates, **eliminating 20+ hours/week** of repeated manual work across the team.
- Refactored legacy reporting scripts into parameterised, documented Python modules, reducing failure rates and enabling handover to other analysts.

Titan Company Limited — Bangalore, India

Data Analyst Intern

Dec 2023 – Jun 2024

- Built **8+ Tableau dashboards** for HR leadership covering hiring funnel, offer acceptance and workforce planning across **500+ employees**, adopted as the standing input to monthly leadership reviews.
- Engineered a **PyQt6 data validation application** that cut manual validation effort from 4 hours/week to 30 minutes and accelerated data parsing by **35%**.

Key Projects

A/B Test Design — Collections Early-Bounce Intervention — *Python, statsmodels, scipy*

Designed a controlled experiment to measure whether an early-bounce alerting intervention — live to 100% of accounts with no holdout — actually changes repayment behaviour. Power analysis showed the intuitive cure-rate metric would need **23 months** to detect the effect that justified the intervention's cost; switching the primary metric to a continuous recovery ratio brought a conclusive test **inside one quarter** on the same population. Implemented stratified randomisation with rerandomisation for covariate balance, quantified spillover from a dealer-level treatment on a customer-level assignment, and built the pre-registered analysis pipeline (Welch t-test with covariate adjustment, Wilson intervals, guardrail checks, BH-corrected segment cuts) with an explicit decision rule.

Operational Reporting Automation Engine — *Python, PostgreSQL, openpyxl*

Multi-sheet Excel automation system tracking asset performance across active, inactive and surrendered statuses, with risk categorisation and audit-ready output. Achieved 50% time reduction at 99% output accuracy.

Cloud-Native E-commerce Decision Engine — *AWS S3, Python, Scikit-Learn, Boto3, NLP*

End-to-end ML pipeline: cloud ingestion, data cleaning, NLP sentiment analysis, Random Forest prediction and rule-based mapping to automated business recommendations.

Education

M.Sc. in Data Science — CHRIST (Deemed to be University), Online

Expected 2027

B.Tech, Computer Science & Engineering (Data Science Specialisation) — CHRIST
(Deemed to be University), Bangalore

Graduated 2024

Certifications

Tableau Desktop Specialist | SQL for Data Science – IBM | Introduction to Data Science – IBM | HackerRank SQL (Advanced)